
title: Models notebook: Distributed Systems layout: note date: 2020-08-12 tags: ...

Models

Models are used to provide abstract, simplified, consistent description of some aspect of interest of distributed system design.

- **physical:** describe types of computers/devices that constitute a system and their interconnectivity without details of specific computer/networking technologies
 - most explicit
- **architectural:** describe system in terms of computational and communication tasks performed by computational elements
 - e.g. client-server, peer-to-peer
- **fundamental:** abstract perspective to describe solutions to particular issues faced by most distributed systems
 - e.g. interaction, failure, security

Physical Models

- **baseline:** extensible set of computer nodes interconnected by computer network for passing of messages
- 3 generations of distributed systems

1970s-80s, early distributed systems

- 10-100 nodes
- local area network, usually Ethernet
- limited Internet connectivity: file transfer, email
- shared local printers and file servers

1990s, Internet-scale distributed systems

- large-scale distributed systems emerged with rapid growth in Internet
- extensible set of nodes interconnected by network of networks (i.e. the Internet)
- significant heterogeneity: networks, computer architecture, operating systems, languages, ...
- emphasis on open standards

Contemporary distributed systems

- nodes in earlier generations were primarily desktop computers which were:
 - *static*: not moving around
 - *discrete*: not embedded in other objects
 - *autonomous*: largely independent of other computers
- in contrast, modern distributed systems don't satisfy these properties:
 - *mobile computing*: nodes have varying location, needing service discovery and spontaneous interoperation
 - *ubiquitous computing*: computers embedded in everyday objects
 - *cloud computing*: pools of nodes together providing a service

Architectural Models

- system architecture: structure in terms of separately specified components and their interrelationships
 - goal: ensure the structure meets present/future demand
 - concerns: reliability, managability, adaptability, cost-effectiveness
- **architectural elements**: interacting components of system
- **architectural patterns**:
- **middleware**

Architectural Elements

Building blocks:

- communicating entities: e.g. threads/processes, nodes
- communication paradigm: e.g. message queue, publish/subscribe

-
- roles and responsibilities: e.g. client, server, peer
 - placement: mapping onto physical distributed infrastructure

Communicating Entities

System-oriented perspective

- **processes:** usually processes are the communicating entities
 - **threads:** strictly threads may be the endpoints of communication
- **nodes:** sensor networks, OS may not support process abstraction

Problem-oriented perspective

- **objects:** objects accessed via interfaces
- **components:** specify interfaces and make dependencies explicit, providing a more complete contract with which to construct the system than objects
- **web services:** closely related to objects/components. Intrinsically integrated with the WWW, using web standards to represent and discover services.
 - software application identified by a URI with interfaces defined, described, discovered as XML
 - supports direct interaction with other software agents via XML message exchange through IP
- objects/components are usually internal to an organisation for tightly coupled applications
- web services are complete services

Communication Paradigms

Direct: coupled senders/receivers

- **interprocess communication:** low-level support for communication between processes in distributed systems
 - message passing primitives, socket programming, multicast
- **remote invocation:** most common communication paradigm for distributed systems
 - 2 way exchange between communicating entities which resulting in remote operation being called
- **request-reply protocol:** pattern on message-passing to support client-server computing

-
- pairwise message exchange
 - most DS use RPC/RMI, but both are supported by underlying request-reply exchanges
 - **remote procedure call (RPC)**: procedures in processes on remote computers can be called as if they are procedures in the local address space.
 - baked in access and location transparency
 - **remote method invocation (RMI)**: resembles RPC but in a world of distributed objects
 - a calling object invokes a method in remote object

Indirect: allow decoupling of senders/receivers Uncoupling

- **space uncoupling**: senders don't need to know who they are sending to
- **time uncoupling**: senders/receivers don't need to exist at the same time

Techniques

- **group communication**: delivery of messages to set of recipients; one-to-many
 - abstraction of group with an ID, which maintains group membership
 - recipients elect to receive messages by joining a group
 - senders send messages to the group using the group ID
- **publish-subscribe**: large number of producers distributing information to a large number of consumers (with different interests); one-to-many
 - uses intermediary service to ensure efficient routing of information from producers to consumers
- **message queues**: point-to-point service; producer sends messages to a specified queue
 - consumer receives messages from the queue
- **tuple space**: processes can place structured data in a persistent tuple space
 - other processes can read/remove tuples by specifying patterns of interest
 - readers/writers don't need to exist simultaneously
- **distributed shared memory**: abstraction for sharing data between processes that don't share physical memory

Roles and Responsibilities

- **client-server:** client processes interact with individual server processes
 - client processes establish connections
 - server processes listen for incoming connections
 - most important and most widely used architecture
- **peer-to-peer:** all processes play similar role as peers
 - resources of each peer are used, so the system resources scale with the number of users
 - substantially more complex than client-server architecture
 - e.g. BitTorrent

Placement

- how objects/services map onto physical distributed infrastructure
- crucial determinant of DS properties: performance, reliability, security
- where to place a given client/server in terms of machines/processes
- needs to account for e.g.:
 - communication pattern between entities
 - reliability of machines and current loading
 - quality of communications
- **mapping services to multiple servers:** services can be implemented as multiple server processes in separate host computers which interact to provide a service to clients
 - can partition/replicate across multiple hosts
 - **cluster:** thousands of commodity processing boards
- **cache:** store of recently used objects that is closer to some clients than objects themselves
 - reduce network traffic and server load
 - may improve performance for client
- **proxy server:** increase availability/performance by reducing load on wide area network
- **mobile code:** applets/Javascript
 - good interactivity, allows asynchronous behaviour (AJAX)
 - server provides code to browser which client browser runs to access services
 - potential security threat

-
- **mobile agent:** running program (code + data) that travels from computer to computer, carrying out a task on someones behalf, typically collecting data, eventually returning with results
 - reduced communication cost and time by replacing remote invocations with local ones
 - big security threat, and web crawlers can still access resources successfully through remote invocations
 - used to install/maintain software on computers within organisation

Architectural Patterns

- patterns build on primitive architectural elements, providing composite recurring structures that work well in particular circumstances
- **layering:** partition system into layers, with a given layer making use of services provided by the layer below. Higher layers are unaware of lower layer implementation details
 - vertical organisation of services into service layers
 - platform: lowest-level hardware/software; e.g. Intel x86/Linux
 - middleware: software that masks heterogeneity and provides useful programming model. Processes interact to implement communication/resource-sharing. Provides building blocks for constructing software
- **tiered architecture:** complements layering. Horizontal partitioning within a layer, separating functionality into different servers. e.g.
 - presentation logic: user interaction/view as presented to user
 - application logic: app-specific processing/business logic
 - data logic: persisted storage; DBMS
 - 3-tier: separates each logical element into a distinct server
- **thin client:** software layer supporting window-based UI local to the user while accessing services on a remote computer
 - allows simple, low-cost devices to be used with a wide range of services
 - poor performance for highly interactive graphical activities: CAD, image processing
 - **Virtual Network Computing (VNC):** remote access to GUI through VNC client via VNC protocol
- **proxy:** support location transparency in RPC/RMI
 - proxy created in local address space to represent remote object: offers same interface as remote object, meaning application programmer calls on the proxy without knowing about the distributed nature

-
- also used for replication/caching
 - **brokerage**: supports interoperability in complex distributed systems

Fundamental Models

- should only contain essential ingredients to understand/reason about aspect's of system behaviour
- make explicit relevant assumptions
- make generalisations concerning what is possible/impossible given those assumptions

Interaction Models

Message passing between processes produces

- communication: information flow
- coordination: synchronisation, ordering

Failure Models

Define and classifies faults

Security Models

Questions

2.1 Provide three specific and contrasting examples of the increasing levels of heterogeneity experienced in contemporary distributed systems as defined in Section 2.2. page 39

- ubiquitous computing: e.g. smart fridges, mobile phones, tablets, laptops, ... with significant differences in performance, input devices, etc.
- mobile computing: nodes may move from location to location
- cloud computing: pools of nodes that together provide a service

2.2 What problems do you foresee in the direct coupling between communicating entities that is implicit in remote invocation approaches? Consequently, what advantages do you anticipate from a level of decoupling as offered by space and time uncoupling? Note: you might want to revisit this answer after reading Chapters 5 and 6. page 43

-
- idempotency

2.3 Describe and illustrate the client-server architecture of one or more major Internet applications (for example, the Web, email or netnews). page 46

the Web: a client (i.e. a browser) opens TCP connections in order to send HTTP requests. A Web server, listening on port 80, responds with a HTTP response.

2.4 For the applications discussed in Exercise 2.1, what placement strategies are employed in implementing the associated services? page 48